Agent checklist: next sessions

UPDATED 13 NOVEMBER 2026 | NEXT SESSION: 14 NOVEMBER

The seven stages below follow the original plan. Preprocessing is now complete according to the teammate. Mariana will pull the changes tomorrow; compatibility with the agents has not yet been checked.

What was done today

[x] Explained both defense conclusions for case 1 and the orchestrator flow.
[x] Added three orchestrator failure tests and simple English comments.
[x] The full suite passed 25 tests before preparing the new cases.
[x] Prepared two new JSON case drafts only. They have NOT been run, tested or added to the live runner.

1. Understand and run the flow - completed for current code

Mariana understands the independent reports and the stop-on-error behaviour. Revisit only the parts affected by the incoming preprocessing changes. The assistant implements and explains changes; comments stay in simple English.

2. Cases and criteria - two new drafts to review

04_documented_payment.json: a 1,200 USD transfer with fictional agreement and invoice excerpts that state the purpose, account pair and due date. Defense should use the documents, but must not claim verified authenticity, innocence or a confirmed false positive.

05_empty_evidence.json: no transaction, timestamp, score or alert flag. The proposed outcome is rejection before any model call. This differs from case 03, which still contains model output and a timestamp.

[ ] Read and approve the scenarios and criteria.
[ ] Explicitly connect the documents in case 04 to CaseContext; the current builder does not load these additional records.
[ ] Add an offline rejection check for case 05 after agreeing the expected behaviour.
[ ] Keep evaluation text and explanatory notes out of model evidence.

Both files are in agents/reference_cases/. They are preparation material, not ready-to-run additions. CASE_NAMES and the experiment runner still use the original three cases. Do not interpret the earlier 25 passing tests as validation of the new drafts.

3. Prosecution / 4. Defense - development baseline completed

Two final Opus sets met the criteria on the original three cases in the assistant assessment. Preserve these outputs. Wider performance, new cases and independent human assessment remain to be checked.

5. Orchestrator - main failure checks completed

[x] Prosecution failure stops before defense.
[x] Defense failure returns no combined result.
[x] Final validation rejects an invented reference from either role.
[ ] Recheck affected paths after integration changes. Define any new result status before changing the current stop-on-error policy.

Tomorrow: inspect the new delivery

SATURDAY, 14 NOVEMBER | REVIEW FIRST, THEN DECIDE

[ ] 1 - Protect local work, then pull

Before Mariana pulls, inspect git status and preserve local edits, especially the existing prompts.py changes. Keep a deliberate commit or backup; do not discard edits or resolve conflicts blindly. Mariana performs the pull.

Done when: The incoming changes are available and local work has not been lost. The assistant has not pulled them in advance.

[ ] 2 - Review the two new cases together

Read case 04 and its fictional documents, then case 05 and the expected rejection. Explain what each case is meant to check.

Done when: Mariana understands the difference between documented purpose, verified legitimacy and no evidence at all.

[ ] 3 - Inspect the preprocessing changes - stage 6

Review the actual changed files, README and generated schemas. Compare the new delivery with agents/service.py and case_context.py before changing the agent code. Do not assume the old interface still works.

Done when: We can name the changed files and explain which changes affect the agents.

[ ] 4 - Write a concrete integration map

Check row_id, feature/metadata locations, account fields, timestamps and timezone, amounts/currencies, missing values, scorer compatibility and daily scoring inputs. Confirm how related_transactions is supplied and whether history coverage is known.

Done when: Every required agent input has a source or an explicit gap. Existing modelling agreements are reused where they still apply.

Tomorrow ends with a small implementation plan

Choose the smallest required adapter changes, where they belong and how they will be checked. Identify missing artifacts or interface questions for the teammate. A completed preprocessing task is not yet evidence of a successful end-to-end agent review.

No immediate need to regenerate the 39 old outputs. First understand the new data and the two draft cases. Only run checks relevant to the changes we choose to implement.

Files to open

agents/reference_cases/04_documented_payment.json
agents/reference_cases/05_empty_evidence.json
agents/service.py
agents/case_context.py
The preprocessing files actually changed by the pull

Following sessions: integrate, then debate

Session 2 - target 15 November, or next available day

Stage 6: implement the agreed integration changes.
[ ] Adapt only fields or paths that changed; keep responsibilities separated.
[ ] Ensure labels and ground-truth-derived fields are excluded from agent evidence.
[ ] Verify row IDs, whole-day scoring and scorer/feature compatibility.
[ ] Check the context for a real prepared row with a simulated text generator first.
[ ] Add document evidence support for case 04 explicitly and keep case 05 as an offline rejection scenario.
[ ] Run relevant offline checks, then the full suite after code changes.

Done when: the new delivery produces the expected context, references remain traceable, and offline checks pass. If artifacts are missing, record the gap and continue with work that does not depend on them.

Session 3 - target 16 November, after integration checks

Finish stage 6 and prepare the larger evaluation.
[ ] Select a valid prepared row that is actually flagged and review its evidence.
[ ] Run one controlled live end-to-end review when the offline handoff is sound.
[ ] Compare claims with the source data and record failures or missing history.
[ ] Evaluate the documented-purpose case only after its documents reach the agents.
[ ] Plan at least ten identified cases for the later live evaluation, including the original three. Keep unseen examples for a meaningful check.

Case 05 cannot count as a successful live report case: it is designed to be rejected before the API. Do not add an invented flag to force it through.

7. Debate and final delivery - after the foundations hold

[ ] Add one defense reply to specific prosecution arguments.
[ ] Preserve the original facts and references; earlier arguments are not evidence.
[ ] Return initial reports and reply with clear labels and a one-round limit.
[ ] Test success, failure and call limits offline.
[ ] Assess at least ten live-review cases in two series; compare the added value of the reply.
[ ] Record model/code version, support for claims, failures, latency and usage when available.
[ ] Write short reproduction instructions and obtain human assessment.

The order is dependency-based, not a deadline to rush. Independent reports can continue to be evaluated while debate is being designed. Do not mark integration or debate complete just because preprocessing is finished.

When to test and where to resume

No tests or model calls were run while creating the two new drafts. After future case wiring or code changes, run focused checks and then the full suite. Reading JSON files needs no API call.

```
.\.venv\Scripts\python.exe -B -m unittest discover -s agents/tests -v
```

Current outputs and explanation: output/output/. Earlier outputs: archive_experiences/ inside that folder. Runner: agents/reference_cases/run_agent_iteration.py. Its --help option makes no API calls.

Next session starts with: preserve local work, pull the teammate changes, review the two case drafts and inspect the new preprocessing interface. Update this checklist with actual results before the following session.